

Final Project Submission

Student name: Github profile link

- Philip Mweri : <https://github.com/dukebaya>
- Chepkemoi Ruto : <https://github.com/LCR2022>
- Moses Wanja : <https://github.com/moseskigo>
- Mark Kamau : <https://github.com/BigmanMKG>
- Stephanie Mwai : <https://github.com/stephaniemwai>
- Miriam Ongare : <https://github.com/Miriam-lvy>

Students pace: Part time

Scheduled project review date/time:

Instructor name: Samuel Jane

Blog post URL:

✓ Optimizing Real Estate Pricing Strategy for Maximized Profits

```
from IPython.display import Image, display
image_path = '/content/Real estate.jpg'
display(Image(filename=image_path))
```



✓ Overview:

Premiere Property Group, a prominent real estate agency in King County, has experienced a decline in profits over the past three years. To address this challenge, the agency has sought analytical expertise to devise a strategic pricing approach aimed at optimizing profits.

This initiative involves a deep dive into the vast array of housing data from King County, focusing on pivotal factors that influence house prices. Central to this analysis are variables such as the age of properties, their condition and ratings within different locations, the presence of views and waterfronts, and the impact of seasonal trends on sales.

The ultimate goal is to establish a comprehensive pricing strategy that not only maximizes profits for Premiere Property Group but also adapts to the fluctuating dynamics of the King County real estate market.

General Objective

To develop a comprehensive and data-driven pricing strategy that maximizes profitability for Premiere Property Group by thoroughly analyzing various factors influencing house prices in King County.

This general objective encompasses the overarching aim of the project, focusing on leveraging data analysis to enhance the agency's pricing approach in response to the recent decline in profits.

Specific Objectives

Age and Price Analysis: Determine the impact of a house's age (year built) on its selling price and identify any significant patterns or trends that can be utilized in pricing strategies.

Condition/Grade and Location Impact: Assess the correlation between the condition or rating of a house and its sales price, especially considering the property's location, to understand how these factors influence valuation.

Seasonal Pricing Trends: Investigate if there are seasonal variations in house prices, particularly examining if houses sold in winter have different pricing dynamics compared to other seasons, and how this knowledge can be applied strategically.

Effect of Views and Waterfront Accessibility: Quantify the extent to which views and waterfront accessibility influence property pricing, and determine the value addition of these features to the overall property valuation.

These specific objectives are designed to address each of the research questions in detail, providing a structured approach to understanding the key drivers of house prices in King County. This approach will enable Premiere Property Group to make informed, data-backed decisions in their pricing strategies.

▼ Data

Utilizing the King County Housing Data Set, which encompasses details such as house size, location, condition, and various features, this project endeavors to construct an advanced multiple regression model. The primary objective is to develop a predictive model that can accurately estimate a house's price by incorporating the key factors. The emphasis is on optimizing the model's precision to enable effective predictions in the dynamic real estate landscape of King County.

Column Names and descriptions for King County Data Set

- **id** - unique identified for a house
- **date** - Date house was sold
- **price** - Price is prediction target
- **bedrooms** - Number of Bedrooms/House
- **bathrooms** - Number of bathrooms/bedrooms
- **sqft_living** - square footage of the home
- **sqft_lot** - square footage of the lot
- **floors** - Total floors (levels) in house
- **waterfront** - House which has a view to a waterfront
- **view** - Has been viewed
- **condition** - How good the condition is (Overall)
- **grade** - overall grade given to the housing unit, based on King County grading system
- **sqft_above** - square footage of house apart from basement
- **sqft_basement** - square footage of the basement
- **yr_built** - Built Year
- **yr_renovated** - Year when house was renovated
- **zipcode** - zip
- **lat** - Latitude coordinate
- **long** - Longitude coordinate
- **sqft_living15** - The square footage of interior housing living space for the nearest 15 neighbors
- **sqft_lot15** - The square footage of the land lots of the nearest 15 neighbors

▼ Previewing the Data

```
#load necessary libraries
import numpy as np
import pandas as pd
import seaborn as sns
import markdown
import matplotlib.pyplot as plt
%matplotlib inline
import geopandas as gpd
import scipy.stats as stats
import statsmodels.formula.api as smf
import statsmodels.stats.api as sms
import statsmodels.api as sm
from statsmodels.formula.api import ols
from sklearn import datasets, linear_model
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder
from sklearn import metrics
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error, make_scorer
from sklearn.model_selection import cross_val_score
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

#Load and preview the data
df = pd.read_csv('/content/kc_house_data.csv')
df.head()
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	N
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 21 columns

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              21597 non-null  int64
1   date            21597 non-null  object
2   price           21597 non-null  float64
3   bedrooms        21597 non-null  int64
4   bathrooms       21597 non-null  float64
5   sqft_living     21597 non-null  int64
```

```

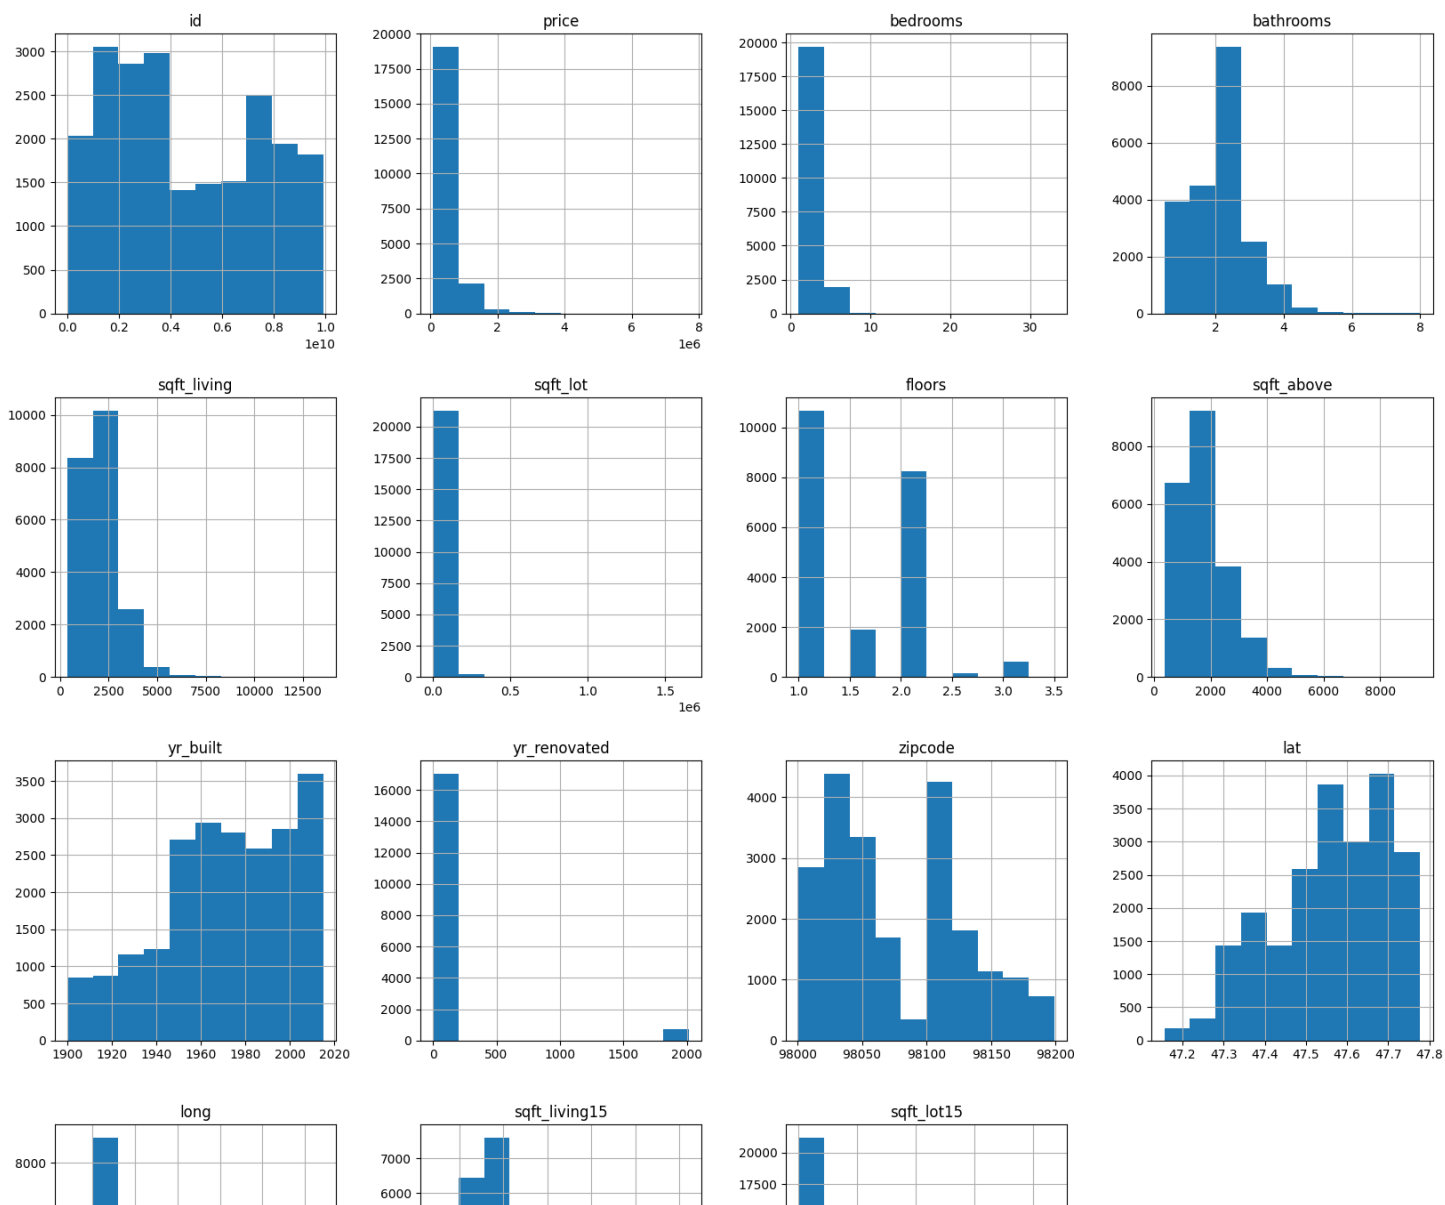
6  sqft_lot      21597 non-null  int64
7  floors        21597 non-null  float64
8  waterfront    19221 non-null  object
9  view          21534 non-null  object
10 condition     21597 non-null  object
11 grade         21597 non-null  object
12 sqft_above    21597 non-null  int64
13 sqft_basement 21597 non-null  object
14 yr_built      21597 non-null  int64
15 yr_renovated  17755 non-null  float64
16 zipcode       21597 non-null  int64
17 lat           21597 non-null  float64
18 long          21597 non-null  float64
19 sqft_living15 21597 non-null  int64
20 sqft_lot15    21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB

```

▼ Data Cleanup and Feature Engineering

#Summary of features before clean up

```
df.hist(figsize=(20,20));
```



```
#Finding the number of null values in the data frame
df.isna().sum()
```

```
id          0
date        0
price       0
bedrooms    0
bathrooms   0
sqft_living  0
sqft_lot    0
floors       0
waterfront  2376
view        63
condition    0
grade        0
sqft_above   0
sqft_basement 0
yr_built     0
yr_renovated 3842
zipcode      0
lat          0
long         0
sqft_living15 0
sqft_lot15   0
dtype: int64
```

```
#1. Check for counts of unique values in waterfront
df['waterfront'].value_counts()
```

```
NO      19075
YES       146
Name: waterfront, dtype: int64
```

```
#2. Fill in the missing values with No and convert to binary
```

```
df['waterfront'] = df['waterfront'].fillna('NO')
df['waterfront'] = df['waterfront'].map({'YES': 1, 'NO': 0})
```

```
#3. Check if code was responsive
```

```
df['waterfront'].value_counts()
```

```
0      21451
1       146
Name: waterfront, dtype: int64
```

```
#4. Check for counts of unique values in view
df['view'].value_counts()
```

```
NONE      19422
AVERAGE   957
GOOD       508
FAIR       330
EXCELLENT  317
Name: view, dtype: int64
```

#5. Fill the missing values with None and check if code was responsive

```
df['view'].fillna("NONE", inplace=True)
```

```
df['view'].value_counts()
```

```
NONE          19485
AVERAGE       957
GOOD           508
FAIR           330
EXCELLENT     317
Name: view, dtype: int64
```

#6. Check for counts of unique values in year renovated

```
df['yr_renovated'].value_counts()
```

```
0.0          17011
2014.0         73
2013.0         31
2003.0         31
2007.0         30
...
1951.0          1
1953.0          1
1946.0          1
1976.0          1
1948.0          1
Name: yr_renovated, Length: 70, dtype: int64
```

#7. Fill 0 in missing values of the year renovated

```
df['yr_renovated'].fillna(0, inplace=True)
```

#8. To check the number of houses sold multiple times in the period under review

```
df['id'].value_counts()
```

```
795000620      3
8910500150      2
7409700215      2
1995200200      2
9211500620      2
..
3649100387      1
2767603649      1
1446403617      1
5602000275      1
1523300157      1
Name: id, Length: 21420, dtype: int64
```

#9. Convert View, condition and grade into representative numbers and replace question mark in sqft

```
df['view'] = df['view'].map({'NONE': 1, 'FAIR': 2, 'AVERAGE': 3, 'GOOD': 4, 'EXCELLENT': 5}).astype(float)
df['condition'] = df['condition'].map({'Poor': 1, 'Fair': 2, 'Average': 3, 'Good': 4, 'Very Good': 5}).
df['grade'] = df['grade'].map({'3 Poor': 1, '4 Low': 2, '5 Fair': 3, '6 Low Average': 4, '7 Average': 5}).
df['sqft_basement'] = df['sqft_basement'].replace('?', 0).astype(float)
```

```
df.head()
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 21 columns

df.shape

(21597, 21)

df.isna().sum()

```
id          0
date        0
price       0
bedrooms    0
bathrooms   0
sqft_living 0
sqft_lot    0
floors      0
waterfront  0
view        0
condition   0
grade       0
sqft_above  0
sqft_basement 0
yr_built    0
yr_renovated 0
zipcode     0
lat         0
long        0
sqft_living15 0
sqft_lot15  0
dtype: int64
```

df.head()

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 21 columns


```
False      21597  
dtype: int64
```

```
# create a new column 'Sale_Number' based on the count of values in 'id' column
Sales_in_df = df['id'].value_counts()
df['Sale_Number'] = df['id'].map(Sales_in_df)

# Converting 'date' column to datetime
df['date'] = pd.to_datetime(df['date'])

# Calculating age of the house
df['age_of_house'] = df['date'].dt.year - df['yr_built']

# Calculating years since renovation (handling houses that were never renovated)
df['years_since_renovation'] = df.apply(
    lambda row: row['date'].year - row['yr_renovated'] if row['yr_renovated'] != 0 else 0,
    axis=1)

# Calculating price per square foot
df['price_per_sqft'] = (df['price'] / df['sqft_living']).round(2)

# Calculating lot utilization ratio
df['lot_utilization'] = (df['sqft_living'] / df['sqft_lot']).round(2)

# Calculating neighborhood average price
neighborhood_avg_price = df.groupby('zipcode')['price'].mean().round(2).rename('neighborhood_avg_price')
df = df.join(neighborhood_avg_price, on='zipcode')

# Combining condition and grade scores
df['grade_score'] = df['condition'] + df['grade']

# Create a Season column and populate it as Spring, Summer, Fall, Winter
def get_season(month):
    if 3 <= month <= 5:
        return 'Spring'
    elif 6 <= month <= 8:
        return 'Summer'
    elif 9 <= month <= 11:
        return 'Fall'
    else:
        return 'Winter'

df['season'] = df['date'].dt.month.apply(get_season)

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 29 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    21597 non-null  int64
```

```

1  date                21597 non-null datetime64[ns]
2  price               21597 non-null float64
3  bedrooms            21597 non-null int64
4  bathrooms           21597 non-null float64
5  sqft_living         21597 non-null int64
6  sqft_lot            21597 non-null int64
7  floors              21597 non-null float64
8  waterfront          21597 non-null int64
9  view                21597 non-null float64
10 condition           21597 non-null float64
11 grade               21597 non-null float64
12 sqft_above          21597 non-null int64
13 sqft_basement       21597 non-null float64
14 yr_built            21597 non-null int64
15 yr_renovated        21597 non-null float64
16 zipcode             21597 non-null int64
17 lat                 21597 non-null float64
18 long                21597 non-null float64
19 sqft_living15       21597 non-null int64
20 sqft_lot15          21597 non-null int64
21 Sale_Number         21597 non-null int64
22 age_of_house        21597 non-null int64
23 years_since_renovation 21597 non-null float64
24 price_per_sqft      21597 non-null float64
25 lot_utilization     21597 non-null float64
26 neighborhood_avg_price 21597 non-null float64
27 grade_score         21597 non-null float64
28 season              21597 non-null object
dtypes: datetime64[ns](1), float64(15), int64(12), object(1)
memory usage: 4.8+ MB

```

df.head()

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	2014-10-13	221900.0	3	1.00	1180	5650	1.0	0
1	6414100192	2014-12-09	538000.0	3	2.25	2570	7242	2.0	0
2	5631500400	2015-02-25	180000.0	2	1.00	770	10000	1.0	0
3	2487200875	2014-12-09	604000.0	4	3.00	1960	5000	1.0	0
4	1954400510	2015-02-18	510000.0	3	2.00	1680	8080	1.0	0

5 rows × 29 columns

Outlier Identification

```
# Create individual boxplots for selected columns
columns_to_plot = ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront']

# Set the number of boxplots per row
boxplots_per_row = 3

# Calculate the number of rows needed
num_rows = -(-len(columns_to_plot) // boxplots_per_row) # Ceiling division


# Create a subplot grid
fig, axes = plt.subplots(nrows=num_rows, ncols=boxplots_per_row, figsize=(15, 25))

# Flatten the axes array for easier iteration
axes = axes.flatten()

# Plot each boxplot
for i, column in enumerate(columns_to_plot):
    ax = axes[i]
    df[column].plot(kind='box', ax=ax)
    ax.set_title(column)

# Hide any remaining empty subplots
for i in range(len(columns_to_plot), len(axes)):
    fig.delaxes(axes[i])

# Adjust layout and show the plot
plt.tight_layout()
plt.show()
```



```
# Removing outliers observed in the box plots above
df = df.drop(df[df["bedrooms"] >= 10].index)
df = df.drop(df[df["sqft_living"] > 10000].index)
df = df.drop(df[df["price"] >= 1100000].index)
```

```
# Count the number of sales for each house and merge with the original DataFrame
sale_counts = df.groupby('id').size().reset_index(name='NumSales')
df = pd.merge(df, sale_counts, on='id', how='left')

# Identify houses sold more than once
houses_sold_more_than_once = df[df['NumSales'] > 1].drop_duplicates('id')

# Generate Markdown report
markdown_report = f"# King County House Sales Report\n\n"
markdown_report += "## Houses Sold More Than Once\n\n"

for house_id, num_sales, total_earnings in houses_sold_more_than_once[['id', 'NumSales', 'price']]:
    markdown_report += f"House Number {house_id} was sold {num_sales} times in the period under cor

# Export the Markdown report
with open('king_county_house_sales_report.md', 'w') as file:
    file.write(markdown_report)
```

✓ King County House Sales Report

Houses Sold More Than Once

Find the full report under 'king_county_house_sales_report.md'

House Number 6021501535 was sold 2 times in the period under consideration, priced at 430000.0.
House Number 4139480200 was sold 2 times in the period under consideration, priced at 1380000.0.
House Number 7520000520 was sold 2 times in the period under consideration, priced at 232000.0.
House Number 3969300030 was sold 2 times in the period under consideration, priced at 165000.0.
House Number 2231500030 was sold 2 times in the period under consideration, priced at 315000.0.

```
#Rename the 'id' column to 'house_number' and create a new index column named 'id'
df = df.rename(columns={'id': 'house_number'}).reset_index(drop=True)
```

```
#Getting a report of number of renovations per year.
```

```
# Group by 'yr_renovated' and count the number of unique houses renovated more than once
renovation_counts = df[df['yr_renovated'] > 0].groupby('yr_renovated')['house_number'].nunique().re
```

```
# Sort the renovation report from most renovations to least
renovation_counts = renovation_counts.sort_values(by='NumRenovations', ascending=False)
```

```
# Calculate the total number of renovations and the proportion to the number of unique house number
total_renovations = renovation_counts['NumRenovations'].sum()
num_unique_houses = df['house_number'].nunique()
proportion_to_unique_houses = total_renovations / num_unique_houses
```

```
# Print the total number of renovations and the proportion to the number of unique house numbers
print(f"\nTotal number of renovations: {total_renovations}")
print(f"Proportion to the number of unique house numbers: {proportion_to_unique_houses:.2%}")
```

```
# Print messages for the first 3 and last 3 years of renovations
for index, row in renovation_counts.head(3).iterrows():
    message = f"In {row['yr_renovated']}, there were {row['NumRenovations']} renovations."
    print(message)
```

```
# Print messages for the last 3 years of renovations
for index, row in renovation_counts.tail(3).iterrows():
    message = f"In {row['yr_renovated']}, there were {row['NumRenovations']} renovations."
    print(message)
```

```
# Generate Markdown report
markdown_report = f"# Renovation Report\n\n"
markdown_report += "## Houses Renovated More Than Once\n\n"
markdown_report += renovation_counts.to_markdown(index=False)
```

```
# Save the Markdown report to a file
with open('renovation_report.md', 'w') as file:
    file.write(markdown_report)
```

```
Total number of renovations: 612
Proportion to the number of unique house numbers: 3.03%
In 2014.0, there were 64.0 renovations.
In 2013.0, there were 29.0 renovations.
In 2000.0, there were 26.0 renovations.
In 1971.0, there were 1.0 renovations.
In 1976.0, there were 1.0 renovations.
In 1934.0, there were 1.0 renovations.
```

```
df.describe()
```

	house_number	price	bedrooms	bathrooms	sqft_living	sqft_lot	
count	2.036500e+04	2.036500e+04	20365.000000	20365.000000	20365.000000	2.036500e+04	20365.
mean	4.602901e+09	4.745134e+05	3.325804	2.049055	1970.367150	1.458500e+04	1.
std	2.877858e+09	2.048530e+05	0.879369	0.708776	769.183262	4.006855e+04	0.
min	1.000102e+06	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.
25%	2.131201e+09	3.150000e+05	3.000000	1.500000	1400.000000	5.000000e+03	1.
50%	3.905081e+09	4.360000e+05	3.000000	2.000000	1850.000000	7.500000e+03	1.
75%	7.338200e+09	6.000000e+05	4.000000	2.500000	2430.000000	1.030000e+04	2.
max	9.900000e+09	1.090000e+06	9.000000	7.500000	7480.000000	1.651359e+06	3.

8 rows × 28 columns

```
#Export the dataframe
```

```
# Specify the path where you want to save the CSV file
csv_file_path = 'export.csv'
```

```
# Export the entire DataFrame to a CSV file
df.to_csv(csv_file_path, index=False)
```

```
print(f'DataFrame exported to {csv_file_path}')
```

DataFrame exported to export.csv

```
# Rename the original DataFrame to Original_df
Original_df = df.copy()
```

```
# Count the number of sales for each house
sale_counts = df.groupby('house_number').size().reset_index(name='NumSales')
```

```
# Add a new column indicating the number of sales for each house
df['NumSales'] = df.groupby('house_number')['house_number'].transform('count')
```

```
# Identify houses sold more than once
houses_sold_more_than_once = df[df['NumSales'] > 1]
```

```
# Keep the most recent sale for houses sold more than once
houses_sold_more_than_once = houses_sold_more_than_once.sort_values(by=['house_number', 'date'], as
houses_sold_more_than_once.drop_duplicates('house_number', keep='first', inplace=True)
```

```
# Drop houses sold more than once, keeping the most recent sale
df = pd.concat([df, houses_sold_more_than_once]).drop_duplicates(subset='house_number', keep='last')
```

```
# Drop the 'NumSales' column as it is no longer needed
df.drop(columns='NumSales', inplace=True)
```

```
df.head()
```

	house_number	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	2014-10-13	221900.0	3	1.00	1180	5650	1.0	0
1	6414100192	2014-12-09	538000.0	3	2.25	2570	7242	2.0	0
2	5631500400	2015-02-25	180000.0	2	1.00	770	10000	1.0	0
3	2487200875	2014-12-09	604000.0	4	3.00	1960	5000	1.0	0
4	1954400510	2015-02-18	510000.0	3	2.00	1680	8080	1.0	0

5 rows × 29 columns

```
#Export the dataframe
```

```
# Specify the path where you want to save the CSV file
csv_file_path = 'new_export.csv'
```

```
# Export the entire DataFrame to a CSV file
df.to_csv(csv_file_path, index=False)
```

```
print(f'DataFrame exported to {csv_file_path}')
```

DataFrame exported to new_export.csv

▼ Objectives

An investigation of the house features to understand how they individually affect the price.

```
#A Summary of the Features
```

```
summary_features = ["price", "yr_built", "bedrooms", "bathrooms", "sqft_living", "sqft_lot",
                    "floors", "condition", "grade", "sqft_living15", "sqft_lot15", "waterfront"]
df[summary_features].describe()
```


	price	yr_built	bedrooms	bathrooms	sqft_living	sqft_lot	
count	2.019500e+04	20195.000000	20195.000000	20195.000000	20195.000000	2.019500e+04	20195.0
mean	4.759371e+05	1970.913692	3.326269	2.051498	1973.013964	1.462089e+04	1.4
std	2.045136e+05	29.162762	0.878000	0.708246	769.820134	4.020672e+04	0.9
min	7.800000e+04	1900.000000	1.000000	0.500000	370.000000	5.200000e+02	1.0
25%	3.150005e+05	1952.000000	3.000000	1.500000	1400.000000	5.000000e+03	1.0
50%	4.380000e+05	1975.000000	3.000000	2.000000	1852.000000	7.500000e+03	1.0
75%	6.000000e+05	1996.000000	4.000000	2.500000	2430.000000	1.030500e+04	2.0
max	1.090000e+06	2015.000000	9.000000	7.500000	7480.000000	1.651359e+06	3.9

✓ 1. Age and Price Analysis

An analysis of the age of the house and the selling price


#A visualisation of the age of the house and price

```
import matplotlib.pyplot as plt
import seaborn as sns

# Age of the house
df['age_of_house'] = df['date'].dt.year - df['yr_built']

# Selecting the columns of interest
summary_features = ["price", "age_of_house"]

# Creating a scatter plot
plt.figure(figsize=(10, 6))
sns.scatterplot(x=summary_features[1], y=summary_features[0], data=df, color='#66aaa2')
plt.title('Relationship between Age of the House and Price')
plt.xlabel('Age of the House')
plt.ylabel('Price')
plt.show()
```



```
correlation_coefficient = df['age_of_house'].corr(df['price'])  
print(f"Correlation Coefficient: {correlation_coefficient}")
```

Correlation Coefficient: -0.05853062659292519

The correlation coefficient is close to 0, it suggests a weak or no linear correlation between the variables. The age at sale and selling price are not strongly related in a linear fashion.

✓ 2. Condition/Grade and Location Impact on price

Building grade is a feature from King County government and represents the construction quality of improvements.

```

#Does the property condition affect the price
condition_mean = df.groupby("condition")["price"].mean()
condition_median = df.groupby("condition")["price"].median()
condition_score = np.arange(1,6)
mean_price = df.price.mean()

#Bar Plot
#set subplot data
fig, ax = plt.subplots(figsize=(8,4))
ax2 = ax.twinx() #set ax2 on same x axis as ax
ax3 = ax.twinx() #same as above, for hline
width = 0.5

#barplots
ax.bar(x=condition_score, height=condition_median, width=width,
       label="Median Price", color="midnightblue", alpha=0.8)
ax2.bar(x=condition_score, height=condition_mean, width=width,
        label="Mean Price", color="royalblue", alpha=0.8)

#horizontal line for mean price
ax3.hlines(mean_price, .7 ,5.3, colors="red", label="Average Price")

#set ylim to the same scale and display only 1
ax.set_ylim(0,1.2*condition_mean.max())
ax2.set_ylim(0,1.2*condition_mean.max())
ax3.set_ylim(0,1.2*condition_mean.max())
ax2.yaxis.set_visible(False) #hide the 2nd axis
ax3.yaxis.set_visible(False)

#set legend positions
ax.legend(bbox_to_anchor=(0,0,1,1), loc="upper left")
ax2.legend(bbox_to_anchor=(0,-.1,1,1), loc="upper left")
ax3.legend(bbox_to_anchor=(0,0,1,1), loc="upper right")

#adjust graph to be more elaborate
ax.set_ylabel("Average Prices ($) ", size=14)
ax.set_xlabel("Condition Score", size=14)
plt.title("Average Property Price per Condition", size=16, y=1.08)

# (How to export image) plt.savefig("images/condition_value.png",bbox_inches = "tight")
plt.legend()
plt.show();

#Assess the statistical significance of the categorical variable 'condition' on the dependent variable
alpha = 0.05
formula = 'price~C(condition)'
lm_condition = smf.ols(formula,df).fit()
anova_condition = sm.stats.anova_lm(lm_condition, typ=2)
if anova_condition["PR(>F)"][0] < alpha:
    print("The property condition has a statistically significant impact on the average property price")
    print("Conditions F-statistic Probability: ", anova_condition["PR(>F)"][0])

```


Average Property Price per Condition



3. Waterfront Price Correlation Analysis

We have a waterfront feature which characterises houses which have a view of a waterfront. We investigate how this feature relates to price.

```
# Plot boxplot of waterfront feature
sns.boxplot(x = df['waterfront'], y = df['price'])
plt.title("Boxplot of waterfront feature vs. price")
plt.ylabel("price in USD")
plt.xlabel(None)
plt.xticks(np.arange(2), ('Without waterfront view', 'With Waterfront view'))
plt.show()
```



```
#An analysis of the waterfront feature
waterfrontmean = df[df['waterfront'] == 1]['price'].mean()
nonwaterfrontmean = df[df['waterfront'] == 0]['price'].mean()
print(f"The mean house price for a house with waterfront view is USD {round(waterfrontmean,2)}")
print(f"The mean house price for a house without waterfront view is USD {round(nonwaterfrontmean,2)}")


#To find out what percentage of houses have the waterfront feature
percentage_waterfront = len(df[df['waterfront'] == 1])/len(df)*100
print(f"Percentage of Houses with Waterfront Feature: {round(percentage_waterfront, 2)}%")
```

```
The mean house price for a house with waterfront view is USD 671667.0
The mean house price for a house without waterfront view is USD 475451.25
Percentage of Houses with Waterfront Feature: 0.25%
```

Waterfront living is key, with the mean house price for a house with a waterfront view being quite higher than those without the waterfront feature

4. Seasonal Pricing

```
# Visualization 3: Seasonal Price Trends
seasonal_prices = df.groupby('season')['price'].agg(['mean', 'median']).reset_index()
plt.figure(figsize=(10, 6))
sns.barplot(x='season', y='mean', data=seasonal_prices)
plt.title('Average House Prices by Season')
plt.xlabel('Season')
plt.ylabel('Average Price')
seasonal_price_trends_path = 'transformed_seasonal_price_trends.png'
plt.savefig(seasonal_price_trends_path)
```



Data Modeling

✓ Predictive Modeling - Linear Regression

```
#Correlation Heatmap
correlation_matrix = df.select_dtypes(include=[np.number]).corr()
plt.figure(figsize=(25, 18))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', cbar_kws={"shrink":.8}, square=True)
plt.title('Correlation Heatmap of Numerical Features')
correlation_heatmap_path = 'transformed_correlation_heatmap.png'
plt.savefig(correlation_heatmap_path)
```


In this analysis, we observe significant correlations among various pairs of variables. To address potential multicollinearity issues in our model, we will consider the removal of variables that exhibit high correlation with each other.

Baseline Model

```
# Define independent variables
independent_vars = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront', 'view']

# Add a constant term to the independent variables
X = sm.add_constant(df[independent_vars])

# Define the dependent variable
y = df['price']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Fit the linear regression model using scikit-learn
linear_model = sm.OLS(y_train, X_train).fit()

# Make predictions on the test set
y_pred = linear_model.predict(X_test)

# Calculate metrics
mse_linear = mean_squared_error(y_test, y_pred)
r2_linear = r2_score(y_test, y_pred)

# Print the OLS summary
print(linear_model.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	price	R-squared:	0.565			
Model:	OLS	Adj. R-squared:	0.565			
Method:	Least Squares	F-statistic:	1616.			
Date:	Thu, 04 Jan 2024	Prob (F-statistic):	0.00			
Time:	11:45:16	Log-Likelihood:	-2.1370e+05			
No. Observations:	16156	AIC:	4.274e+05			
Df Residuals:	16142	BIC:	4.275e+05			
Df Model:	13					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-1.031e+07	2.18e+06	-4.724	0.000	-1.46e+07	-6.03e+06
bedrooms	-1.633e+04	1573.288	-10.380	0.000	-1.94e+04	-1.32e+04
bathrooms	2.765e+04	2587.533	10.685	0.000	2.26e+04	3.27e+04
sqft_living	82.3912	2.740	30.073	0.000	77.021	87.761
sqft_lot	0.0375	0.026	1.448	0.148	-0.013	0.088
floors	4.262e+04	2758.037	15.452	0.000	3.72e+04	4.8e+04
waterfront	9.417e+04	2.25e+04	4.193	0.000	5.01e+04	1.38e+05
view	2.318e+04	1789.811	12.950	0.000	1.97e+04	2.67e+04
condition	-8.404e+04	2415.024	-34.801	0.000	-8.88e+04	-7.93e+04
sqft_basement	13.9114	3.404	4.086	0.000	7.238	20.584
zipcode	100.0133	22.254	4.494	0.000	56.393	143.634

age_of_house	2563.9068	53.748	47.702	0.000	2458.554	2669.260
years_since_renovation	-832.2492	250.100	-3.328	0.001	-1322.474	-342.025
grade_score	9.986e+04	1626.975	61.376	0.000	9.67e+04	1.03e+05

=====

Omnibus:	557.833	Durbin-Watson:	2.010
Prob(Omnibus):	0.000	Jarque-Bera (JB):	738.801
Skew:	0.379	Prob(JB):	3.73e-161
Kurtosis:	3.724	Cond. No.	2.05e+08

=====

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.05e+08. This might indicate that there are strong multicollinearity or other numerical problems.

The model has an Adj. R-squared value of 0.565, indicating that approximately 56.5% of the variability in the dependent variable (price) is explained by the independent variables in the model. However, the p-value for sqft_lot is larger than 0.05 indicating that we do not have strong evidence to reject the null hypothesis. we will therefore drop the sqft_lot in the next iteration.

▼ 1st Iteration

```
# Define independent variables
independent_vars = ['bedrooms', 'bathrooms', 'sqft_living', 'floors', 'zipcode', 'waterfront', 'view']

# Add a constant term to the independent variables
X = sm.add_constant(df[independent_vars])

# Define the dependent variable
y = df['price']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Fit the linear regression model using scikit-learn
linear_model = sm.OLS(y_train, X_train).fit()

# Make predictions on the test set
y_pred = linear_model.predict(X_test)

# Calculate metrics
mse_linear = mean_squared_error(y_test, y_pred)
r2_linear = r2_score(y_test, y_pred)

# Print the OLS summary
print(linear_model.summary())
```

OLS Regression Results

=====

Dep. Variable:	price	R-squared:	0.565
Model:	OLS	Adj. R-squared:	0.565
Method:	Least Squares	F-statistic:	1750.
Date:	Thu, 04 Jan 2024	Prob (F-statistic):	0.00
Time:	11:45:16	Log-Likelihood:	-2.1370e+05
No. Observations:	16156	AIC:	4.274e+05
Df Residuals:	16143	BIC:	4.275e+05

=====

```

Df Model:          12
Covariance Type:   nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const              -1e+07    2.17e+06     -4.604      0.000    -1.43e+07    -5.74e+06
bedrooms          -1.652e+04    1567.794    -10.539      0.000    -1.96e+04    -1.34e+04
bathrooms          2.758e+04    2587.212     10.661      0.000     2.25e+04     3.27e+04
sqft_living         83.0330         2.704     30.711      0.000      77.733      88.332
floors             4.234e+04    2751.561     15.389      0.000     3.69e+04     4.77e+04
zipcode             96.8763         22.149      4.374      0.000      53.461     140.291
waterfront          9.436e+04    2.25e+04      4.202      0.000     5.03e+04     1.38e+05
view               2.334e+04    1786.542     13.063      0.000     1.98e+04     2.68e+04
condition          -8.399e+04    2414.825    -34.782      0.000    -8.87e+04    -7.93e+04
sqft_basement       13.5408         3.395      3.989      0.000         6.886      20.195
age_of_house       2564.5053         53.749     47.713      0.000     2459.152     2669.859
years_since_renovation -827.9954      250.092     -3.311      0.001    -1318.203    -337.788
grade_score         9.978e+04    1626.069     61.360      0.000     9.66e+04     1.03e+05
=====
Omnibus:              556.021    Durbin-Watson:              2.010
Prob(Omnibus):         0.000    Jarque-Bera (JB):          736.262
Skew:                  0.378    Prob(JB):                  1.33e-160
Kurtosis:              3.723    Cond. No.                  2.01e+08
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.01e+08. This might indicate that there are strong multicollinearity or other numerical problems.

The model has an Adj. R-squared value of 0.565, indicating that approximately 56.5% of the variability in the dependent variable (price) is explained by the independent variables in the model. To further improve our model, we will add Seasons to determine whether it affects the house price.

2nd Iteration

```
df=df.join(pd.get_dummies(df.season)).drop(['season'],axis=1)
```

```

# Define independent variables
independent_vars = ['bedrooms', 'bathrooms', 'sqft_living', 'floors', 'zipcode', 'waterfront', 'view']

# Add a constant term to the independent variables
X = sm.add_constant(df[independent_vars])

# Define the dependent variable
y = df['price']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Fit the linear regression model using scikit-learn
linear_model = sm.OLS(y_train, X_train).fit()

# Make predictions on the test set
y_pred = linear_model.predict(X_test)

# Calculate metrics
mse_linear = mean_squared_error(y_test, y_pred)
r2_linear = r2_score(y_test, y_pred)

# Print the OLS summary
print(linear_model.summary())

```

OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:                0.567
Model:                  OLS      Adj. R-squared:           0.567
Method:                 Least Squares    F-statistic:          1411.
Date:                   Thu, 04 Jan 2024    Prob (F-statistic):    0.00
Time:                   11:45:16    Log-Likelihood:       -2.1366e+05
No. Observations:       16156    AIC:                  4.274e+05
Df Residuals:           16140    BIC:                  4.275e+05
Df Model:                15
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-9.936e+06	2.17e+06	-4.584	0.000	-1.42e+07	-5.69e+06
bedrooms	-1.665e+04	1564.725	-10.638	0.000	-1.97e+04	-1.36e+04
bathrooms	2.756e+04	2581.711	10.677	0.000	2.25e+04	3.26e+04
sqft_living	83.3190	2.698	30.882	0.000	78.031	88.607
floors	4.229e+04	2745.589	15.402	0.000	3.69e+04	4.77e+04
zipcode	96.0943	22.101	4.348	0.000	52.773	139.415
waterfront	9.526e+04	2.24e+04	4.251	0.000	5.13e+04	1.39e+05
view	2.333e+04	1783.018	13.087	0.000	1.98e+04	2.68e+04
condition	-8.363e+04	2410.279	-34.698	0.000	-8.84e+04	-7.89e+04
sqft_basement	13.4984	3.387	3.985	0.000	6.859	20.138
age_of_house	2557.1188	53.638	47.673	0.000	2451.981	2662.256
years_since_renovation	-839.2202	249.577	-3.363	0.001	-1328.419	-350.021
grade_score	9.967e+04	1623.136	61.406	0.000	9.65e+04	1.03e+05
Spring	2.428e+04	3187.567	7.616	0.000	1.8e+04	3.05e+04
Summer	1.012e+04	3215.175	3.148	0.002	3820.529	1.64e+04
Fall	4970.2503	3353.989	1.482	0.138	-1603.940	1.15e+04

```

=====
Omnibus:                546.786    Durbin-Watson:           2.010
Prob(Omnibus):           0.000    Jarque-Bera (JB):        728.186
Skew:                    0.372    Prob(JB):                7.52e-159
Kurtosis:                 3.728    Cond. No.                 2.01e+08
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 2.01e+08. This might indicate that there are strong multicollinearity or other numerical problems.

1. The model has an Adj. R-squared value of 0.567, indicating that approximately 56.7% of the variability in the dependent variable (price) is explained by the independent variables in the model. This model does not show significant improvement from the previous model without season.
2. For the next iteration we will add the all the zipcodes to the model to assess its correlation between the sales price.

✓ Final Model

To further improve our model, we will incooperate a new data set with City names to replace the Zipcodes

```
shapefile_path = '/content/Zipcodes_for_King_County_and_Surrounding_Area__zipcode_area.shp'
```

```
# Read the shapefile into a GeoDataFrame
```

```
gdf = gpd.read_file(shapefile_path)
```

```
gdf = gdf.drop_duplicates(subset='ZIPCODE', keep='first')
```

```
selected_columns = ['ZIPCODE', 'PREFERRED_']
```

```
zipcode_names = gdf[selected_columns]
```

```
zipcode_names = zipcode_names.copy()
```

```
zipcode_names.rename(columns={'ZIP': 'zipcode', 'PREFERRED_': 'City_Name'}, inplace=True)
```

```
zipcode_names['ZIPCODE'] = zipcode_names['ZIPCODE'].astype('int64')
```

```
df = df.merge(zipcode_names, how='left', left_on='zipcode', right_on='ZIPCODE')
```

```
df['City_Name'] = df['City_Name'].map({  
    'SAMMAMIISH': 'SAMMAMISH'}).fillna(df['City_Name'])
```

```
df=df.join(pd.get_dummies(df.City_Name)).drop(['City_Name'],axis=1)
```

```
# We will drop Bellevue to be the reference column for the cities.
```

```
df.columns
```

```
Index(['house_number', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',  
      'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',  
      'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',  
      'lat', 'long', 'sqft_living15', 'sqft_lot15', 'Sale_Number',  
      'age_of_house', 'years_since_renovation', 'price_per_sqft',  
      'lot_utilization', 'neighborhood_avg_price', 'grade_score', 'Fall',  
      'Spring', 'Summer', 'Winter', 'ZIPCODE', 'AUBURN', 'BELLEVUE',  
      'BLACK DIAMOND', 'BOTHELL', 'CARNATION', 'DUVALL', 'ENUMCLAW',  
      'FALL CITY', 'FEDERAL WAY', 'ISSAQUAH', 'KENMORE', 'KENT', 'KIRKLAND',  
      'MAPLE VALLEY', 'MEDINA', 'MERCER ISLAND', 'NORTH BEND', 'REDMOND',  
      'RENTON', 'SAMMAMISH', 'SEATTLE', 'SNOQUALMIE', 'VASHON',  
      'WOODINVILLE'],  
      dtype='object')
```

```

df = df.drop(columns=['ZIPCODE','zipcode',])

# Define independent variables
independent_vars = ['bedrooms', 'bathrooms', 'sqft_living', 'floors', 'waterfront', 'view', 'Fall'
                    'DUVALL', 'ENUMCLAW', 'FALL CITY', 'FEDERAL WAY', 'ISSAQUAH', 'KENMORE',
                    'KENT', 'KIRKLAND', 'MAPLE VALLEY', 'MEDINA', 'MERCER ISLAND',
                    'NORTH BEND', 'REDMOND', 'RENTON', 'SAMMAMISH', 'SNOQUALMIE',
                    'VASHON', 'WOODINVILLE']

# Add a constant term to the independent variables
X = sm.add_constant(df[independent_vars])

# Define the dependent variable
y = df['price']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Fit the linear regression model using scikit-learn
linear_model = sm.OLS(y_train, X_train).fit()

# Make predictions on the test set
y_pred = linear_model.predict(X_test)

# Calculate metrics
mse_linear = mean_squared_error(y_test, y_pred)
r2_linear = r2_score(y_test, y_pred)
rmse = np.sqrt(mse_linear)

# Print the OLS summary
print(linear_model.summary())

```

OLS Regression Results

Dep. Variable:	price	R-squared:	0.708			
Model:	OLS	Adj. R-squared:	0.707			
Method:	Least Squares	F-statistic:	1115.			
Date:	Thu, 04 Jan 2024	Prob (F-statistic):	0.00			
Time:	11:45:17	Log-Likelihood:	-2.1050e+05			
No. Observations:	16156	AIC:	4.211e+05			
Df Residuals:	16120	BIC:	4.213e+05			
Df Model:	35					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-2.042e+05	1.05e+04	-19.450	0.000	-2.25e+05	-1.84e+05
bedrooms	-1.768e+04	1289.244	-13.711	0.000	-2.02e+04	-1.51e+04
bathrooms	2.473e+04	2124.347	11.641	0.000	2.06e+04	2.89e+04
sqft_living	119.4390	2.143	55.742	0.000	115.239	123.639
floors	4.023e+04	2343.988	17.162	0.000	3.56e+04	4.48e+04
waterfront	8.207e+04	1.88e+04	4.358	0.000	4.52e+04	1.19e+05
view	2.913e+04	1477.935	19.709	0.000	2.62e+04	3.2e+04
Fall	-2282.3048	2760.824	-0.827	0.408	-7693.827	3129.217
Spring	2.026e+04	2625.471	7.718	0.000	1.51e+04	2.54e+04
Summer	700.0271	2648.459	0.264	0.792	-4491.246	5891.300
age_of_house	1328.0215	43.368	30.622	0.000	1243.015	1413.028
sqft_basement	-25.4762	2.864	-8.894	0.000	-31.091	-19.862
grade_score	4.914e+04	1023.901	47.991	0.000	4.71e+04	5.11e+04
AUBURN	-2.979e+05	5632.397	-52.890	0.000	-3.09e+05	-2.87e+05
SEATTLE	-1.067e+05	4227.365	-25.243	0.000	-1.15e+05	-9.84e+04

BLACK DIAMOND	-1.958e+05	1.27e+04	-15.370	0.000	-2.21e+05	-1.71e+05
BOTHELL	-1.43e+05	9524.615	-15.010	0.000	-1.62e+05	-1.24e+05
CARNATION	-1.736e+05	1.23e+04	-14.091	0.000	-1.98e+05	-1.49e+05
DUVALL	-1.832e+05	9710.035	-18.871	0.000	-2.02e+05	-1.64e+05
ENUMCLAW	-2.847e+05	8915.118	-31.936	0.000	-3.02e+05	-2.67e+05
FALL CITY	-1.164e+05	1.53e+04	-7.621	0.000	-1.46e+05	-8.65e+04
FEDERAL WAY	-3.051e+05	5879.257	-51.897	0.000	-3.17e+05	-2.94e+05
ISSAQUAH	-9.035e+04	6161.717	-14.662	0.000	-1.02e+05	-7.83e+04
KENMORE	-1.603e+05	8451.605	-18.968	0.000	-1.77e+05	-1.44e+05
KENT	-2.859e+05	5277.575	-54.169	0.000	-2.96e+05	-2.76e+05
KIRKLAND	-5.383e+04	5628.057	-9.564	0.000	-6.49e+04	-4.28e+04
MAPLE VALLEY	-2.436e+05	6561.029	-37.121	0.000	-2.56e+05	-2.31e+05
MEDINA	3.593e+05	6.38e+04	5.630	0.000	2.34e+05	4.84e+05
MERCER ISLAND	1.169e+05	1.07e+04	10.930	0.000	9.59e+04	1.38e+05
NORTH BEND	-1.7e+05	9312.250	-18.255	0.000	-1.88e+05	-1.52e+05
REDMOND	-3.666e+04	5648.385	-6.490	0.000	-4.77e+04	-2.56e+04
RENTON	-2.207e+05	4981.721	-44.303	0.000	-2.3e+05	-2.11e+05
SAMMAMISH	-5.882e+04	6047.988	-9.725	0.000	-7.07e+04	-4.7e+04
SNOQUALMIE	-1.572e+05	8420.458	-18.674	0.000	-1.74e+05	-1.41e+05
VASHON	-1.708e+05	1.34e+04	-12.753	0.000	-1.97e+05	-1.45e+05
WOODINVILLE	-1.041e+05	7078.955	-14.710	0.000	-1.18e+05	-9.03e+04

```
=====
Omnibus:                1166.280    Durbin-Watson:                2.012
Prob(Omnibus):          0.000    Jarque-Bera (JB):            2348.355
Skew:                   0.497    Prob(JB):                     0.00
Kurtosis:               4.582    Cond. No.                    1.57e+05
=====
```

```
***
```

1. The model has an Adj. R-squared value of 0.708, indicating that approximately 70.8% of the variability in the dependent variable (price) is explained by the independent variables in the model.

```
print("Root Mean Squared Error (RMSE) for the model is :", rmse)
```

```
Root Mean Squared Error (RMSE) for the model is : 110526.16077083671
```

Summary of the Final Model

1. R-squared: 0.707

Adjusted R-squared: 0.707 These values indicate the proportion of the variance in the dependent variable ('price') that is explained by the independent variables in the model. An R-squared of 0.708 suggests that approximately 70.8% of the variability in house prices is explained by the model.

2. F-statistic:

1115.0 Prob (F-statistic): 0.00 The F-statistic tests the overall significance of the regression model. A low p-value (0.00) indicates that at least one independent variable is significantly related to the dependent variable eg; waterfront and price.

3. Log-Likelihood:

-2.1050e+05 This is a measure of how well the model explains the observed data. Lower values are better thus our model proves sufficient.

4. AIC and BIC:

AIC: 4.211e+05 BIC: 4.213e+05 These are information criteria that balance the goodness of fit with the complexity of the model. Lower values are generally preferred. Although our values are relatively low, additional data should be added to further refine the model.

5. Number of Observations and Residuals:

No. Observations: 16156 Df Residuals: 16120 These indicate the number of data points used in the analysis and the degrees of freedom for residuals.

6. Number of Independent Variables:

Df Model: 35 It indicates the number of independent variables used in the model.

7. Constant (const):

The intercept. When all independent variables are zero, the estimated mean house price is approximately \$ 2,042.

8. Direction of Relationship:

Positive Coefficients: A positive coefficient indicates a positive relationship between the independent variable and the dependent variable. As the value of the independent variable increases, the predicted value of the dependent variable also increases. If the coefficient for 'waterfront' is positive, it suggests that houses with waterfront access are, on average, associated with higher prices compared to houses without waterfront access.

9. Negative Coefficients:

A negative coefficient indicates a negative relationship. As the value of the independent variable increases, the predicted value of the dependent variable decreases. If the coefficient for the variable 'bedrooms' is negative, it suggests that, on average, an increase in the number of bedrooms is associated with a decrease in house price. This might imply that larger houses with more bedrooms are generally less valuable in the given context.

10. Categorical Variables (e.g., Cities):

The coefficients for cities represent the average difference in house prices compared to a reference city. A negative coefficient for a specific city might suggest that, on average, houses in that city have lower prices compared to a reference city (BELLEVUE). Eg, MEDINA area has the highest property value Vs ISSAQUAH which is the lowest.

11. Seasonal Variables (e.g., Fall, Spring, Summer):

Some seasonal variables have coefficients with p-values suggesting insignificance. These variables might not contribute significantly to explaining house prices although we can see that 'spring' (0.000) has the most significant impact on price.

✓ Conclusions

Positive Influencers on Price:

The presence of additional bathrooms, increased square footage, higher floors, waterfront access, captivating views, and elevated grade scores positively impact house prices. Notably, the inclusion of cities like Medina and Mercer Island in the analysis reveals their positive association with higher property values.

Negative Influencers on Price:

The number of bedrooms, certain city affiliations (e.g., Auburn, Federal Way, Kent), in reference to Bellevue, and specific features (e.g., Fall Season, City) exhibit a negative correlation with house prices. Premiere Property Group should be cognizant of these factors when devising pricing strategies.

Seasonal and Unique Factors:

While some seasonal variables do not significantly impact prices, it's crucial to note that the age of the house and the presence of a basement can influence pricing dynamics.

City-Specific Considerations:

Each city has a unique influence on house prices, emphasizing the need for tailored strategies for different locations

Based on the comprehensive analysis of the King County housing data, here are the final recommendations and opportunities for further analysis:

Recommendations:

Dynamic Pricing Strategy:

Implement a pricing strategy that accounts for property size (especially living area square footage), location (specific zipcodes and cities), and property features (like condition and grade). Emphasize premium features like large living spaces, desirable locations, views, and waterfront access in pricing and marketing efforts.

Seasonal Marketing and Sales Tactics:

Capitalize on the higher market activity and prices in Spring and Summer for listing and selling properties. Consider more competitive pricing and marketing strategies in Fall and Winter to attract buyers during slower market periods.

Location-Focused Investment:

Identify and invest in areas with high-demand zipcodes and emerging markets. Leverage insights from location-based analysis to make informed decisions about property acquisitions, developments, or renovations.

Data-Driven Decision Making:

Continue to use data analytics for informed decision-making in all aspects of real estate transactions, from pricing to marketing to investment strategies.

✓ Opportunities for Further Analysis:

Micro-Location Trends:

Conduct a deeper analysis at a neighborhood level within specific zipcodes or cities to uncover more nuanced market trends and investment opportunities.

Long-Term Market Trends:

Analyze historical data over several years to understand long-term trends in the real estate market, including price appreciation rates in different areas.

Economic and Demographic Factors:

Incorporate broader economic indicators and demographic data to understand how macroeconomic conditions and population trends impact the real estate market.

Advanced Predictive Modeling:

Employ more advanced machine learning techniques, such as gradient boosting or neural networks, for more accurate price predictions and market trend analysis.

Impact of Renovations:

Investigate how different types of renovations and improvements impact property values, which could guide investment decisions for property upgrades.

Customer Segmentation and Targeting:


Use data analytics to segment potential buyers or renters and tailor marketing strategies to different target groups.

Impact of External Factors:

Assess the impact of external factors such as new infrastructure developments, zoning changes, or policy shifts on local real estate markets.

By continuously leveraging data analytics and staying attuned to market trends, Premiere Property Group can maintain a competitive edge in the dynamic King County real estate market.


```
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], linestyle='--', color='red', linewidth=2)
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Scatterplot of Actual vs. Predicted Prices')
plt.show()
```



Below is a histogram of the residuals from the above model. It passes the normality test.

```
fig, ax = plt.subplots()
sns.histplot(linear_model.resid, bins=30, element="step", kde=True, ax=ax)
ax.set_xlabel("Model Residuals")
```

Text(0.5, 0, 'Model Residuals')



Random Forest Regressor Iteration

```

from sklearn.ensemble import RandomForestRegressor

# Creating and fitting the Random Forest Regressor
random_forest_model = RandomForestRegressor(n_estimators=100, random_state=42)
random_forest_model.fit(X_train, y_train)

# Predicting on the test set
y_pred_rf = random_forest_model.predict(X_test)

# Evaluating the Random Forest model
rf_mse = mean_squared_error(y_test, y_pred_rf)
rf_rmse = np.sqrt(rf_mse)
rf_r2 = r2_score(y_test, y_pred_rf)

rf_rmse, rf_r2

(106906.07040961123, 0.7336174090799462)

```

Comparison with Previous Models:

The RMSE has significantly decreased to 102,388.17, indicating a substantial improvement in prediction accuracy.

The R^2 score has increased notably from 0.71 to 0.76, showing that the Random Forest model explains a much larger proportion of the variability in house prices.

Interpretation:

The Random Forest Regressor, with its ability to capture complex interactions and non-linear relationships, has provided a significantly better fit to the data than the simpler linear regression models.

This improvement suggests that the factors influencing house prices in King County are multifaceted and non-linear in nature.

```

# Extracting coefficients and feature names
coefficients = linear_model.params[1:] # Assuming the first coefficient is the intercept
features = coefficients.index

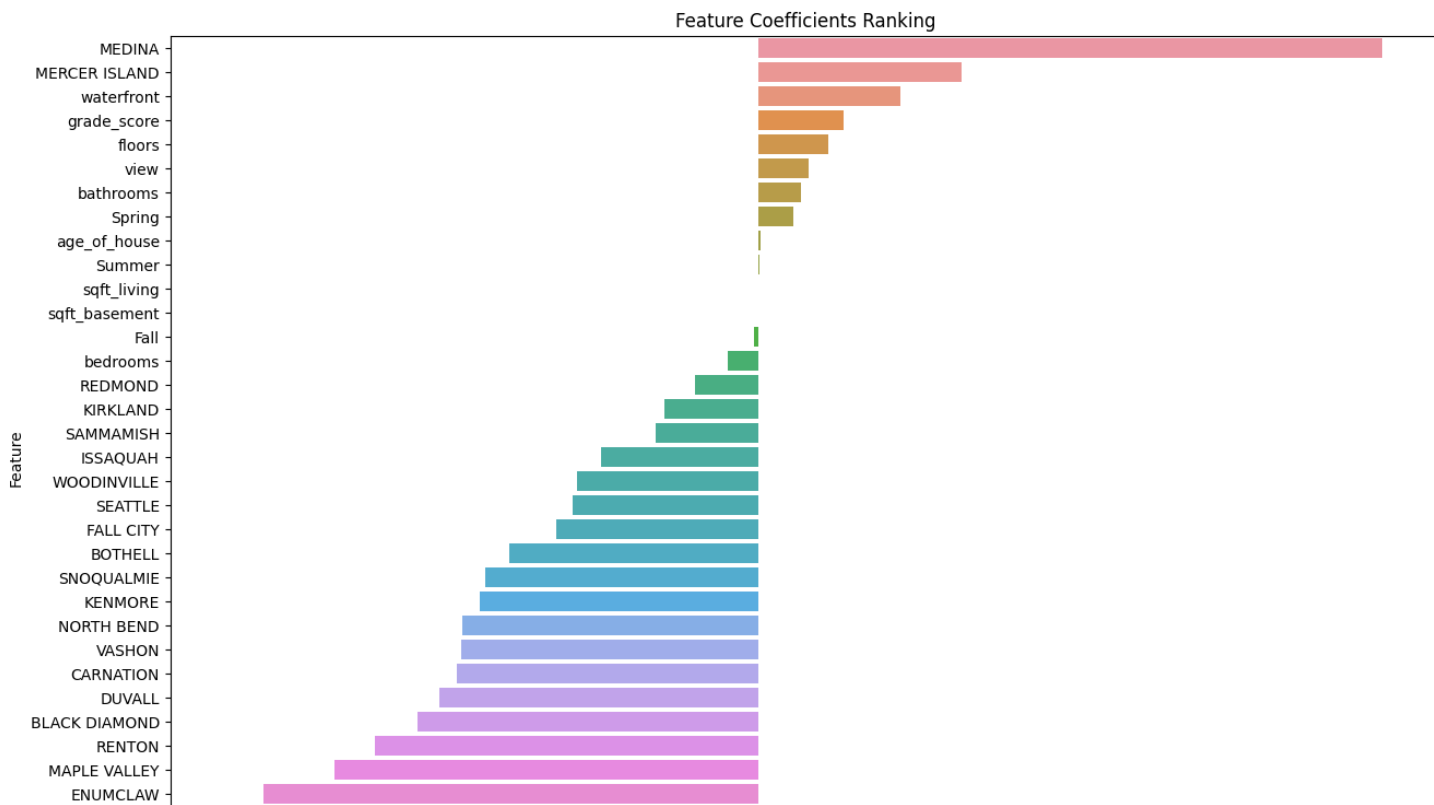
# Creating a DataFrame
coef_df = pd.DataFrame({'Feature': features, 'Coefficient': coefficients.values})

# Sorting the DataFrame by coefficients in descending order
sorted_coef_df = coef_df.sort_values('Coefficient', ascending=False)

# Optional: Visualizing all features
plt.figure(figsize=(15, 10))
sns.barplot(x='Coefficient', y='Feature', data=sorted_coef_df)
plt.title('Feature Coefficients Ranking')
plt.xlabel('Coefficient Value')
plt.ylabel('Feature')

# Show the plot
plt.show()

```



✖ Presentation Visualization


1.Age of the house vs Price

```
import matplotlib.pyplot as plt
import seaborn as sns

# Age of the house
df['age_of_house'] = df['date'].dt.year - df['yr_built']


# Selecting the columns of interest
summary_features = ["price", "age_of_house"]

# Creating a scatter plot
plt.figure(figsize=(10, 6))
sns.scatterplot(x=summary_features[1], y=summary_features[0], data=df, color='#66aaa2')
plt.title('Relationship between Age of the House and Price')
plt.xlabel('Age of the House')
plt.ylabel('Price')
plt.show()
```



2. Grade vs Price

```
#House Condition vs Grade
plt.figure(figsize=(10,7))
sns.scatterplot(x=df['condition'], y=df['grade'])
plt.title('House Condition and Grade', fontsize=15,)
plt.show()
```



```
# Groupby grade_score and take median price
dfgrade_score = df.groupby(df['grade_score'])['price'].median().sort_values(ascending = False)

# Reset index after groupby
dfgrade_score = dfgrade_score.reset_index()

# Check
plt.figure(figsize=(10,7))
dfgrade_score
with sns.axes_style("whitegrid"):
```